


Heapsort



Heaps

A *heap* can be seen as a complete binary tree:




What makes a binary tree complete?

Is the example above complete?



Heaps

A *heap* can be seen as a complete binary tree:




The book calls them “nearly complete” binary trees; can think of unfilled slots as null pointers



Heaps

In practice, heaps are usually implemented as arrays:



Heaps

To represent a complete binary tree as an array:


The root node is $A[1]$

Node i is $A[i]$

The parent of node i is $A[i/2]$ (note: integer divide)

The left child of node i is $A[2i]$

The right child of node i is $A[2i + 1]$



Referencing Heap Elements

So...

```
Parent(A[i]) { return A[ $\lfloor i/2 \rfloor$ ]; }
```

```
Left(A[i]) { return A[2*i]; }
```

```
right(A[i]) { return A[2*i + 1]; }
```



The Heap Property

Heaps also satisfy the *heap property*:

$$A[\textit{Parent}(A[i])] \geq A[i] \quad \text{for all nodes } i > 1$$

In other words, the value of a node is at most the value of its parent

Where is the largest element in a heap stored?



Heap Height

Definitions:

The *height* of a node in the tree = the number of edges on the longest downward path to a leaf

The height of a tree = the height of its root

What is the height of an n -element heap? Why?

This is nice: basic heap operations take at most time proportional to the height of the heap



Heap Operations: Heapify()

Heapify() : maintain the heap property

Given: a node i in the heap with children l and r

Given: two subtrees rooted at l and r , assumed to be heaps

Problem: The subtree rooted at i may violate the heap property (*How?*)

Action: let the value of the parent node “float down” so subtree at i satisfies the heap property

What do you suppose will be the basic operation between i , l , and r ?




Heap Operations: Heapify()

```
Heapify(A, i)
{
    l = Left(i); r = Right(i);
    if (l <= heap_size(A) && A[l] > A[i])
        largest = l;
    else
        largest = i;
    if (r <= heap_size(A) && A[r] > A[largest])
        largest = r;
    if (largest != i)
        Swap(A, i, largest);
    Heapify(A, largest);
}
```



Heapify() Example




$A =$

16	4	10	14	7	9	3	2	8	1
----	---	----	----	---	---	---	---	---	---



Heapify() Example





A =


16	4	10	14	7	9	3	2	8	1
----	---	----	----	---	---	---	---	---	---



Heapify() Example



Heapify() Example




$A =$

16	14	10	4	7	9	3	2	8	1
----	----	----	---	---	---	---	---	---	---



Heapify() Example





A =


16	14	10	4	7	9	3	2	8	1
----	----	----	---	---	---	---	---	---	---



Heapify() Example



Heapify() Example




$A =$

16	14	10	8	7	9	3	2	4	1
----	----	----	---	---	---	---	---	---	---



Heapify() Example




$A =$

16	14	10	8	7	9	3	2	4	1
----	----	----	---	---	---	---	---	---	---



Heapify() Example



$A =$

16	14	10	8	7	9	3	2	4	1
----	----	----	---	---	---	---	---	---	---



Analyzing Heapify(): Informal

*Aside from the recursive call, what is the running time of **Heapify()**?*

*How many times can **Heapify()** recursively call itself?*

*What is the worst-case running time of **Heapify()** on a heap of size n ?*



Analyzing Heapify(): Formal

Fixing up relationships between i , l , and r takes $\Theta(1)$ time

If the heap at i has n elements, how many elements can the subtrees at l or r have?

Draw it

Answer: $2n/3$ (worst case: bottom row 1/2 full)

So time taken by **Heapify()** is given by

$$T(n) \leq T(2n/3) + \Theta(1)$$



Analyzing Heapify(): Formal

So we have

$$T(n) \leq T(2n/3) + \Theta(1)$$

By case 2 of the Master Theorem,

$$T(n) = O(\lg n)$$

Thus, **Heapify()** takes logarithmic time



Heap Operations: BuildHeap()

We can build a heap in a bottom-up manner by running **Heapify()** on successive subarrays

Fact: for array of length n , all elements in range $A[\lfloor n/2 \rfloor + 1 .. n]$ are heaps (*Why?*)

So:

Walk backwards through the array from $n/2$ to 1, calling **Heapify()** on each node.

Order of processing guarantees that the children of node i are heaps when i is processed



BuildHeap()


```
// given an unsorted array A, make A a heap
BuildHeap(A)
{
    heap_size(A) = length(A);
    for (i =  $\lfloor \text{length}[A] / 2 \rfloor$  downto 1)
        Heapify(A, i);
}
```



BuildHeap() Example

Work through example

$A = \{4, 1, 3, 2, 16, 9, 10, 14, 8, 7\}$



Analyzing BuildHeap()

Each call to **Heapify()** takes $O(\lg n)$ time

There are $O(n)$ such calls (specifically, $\lfloor n/2 \rfloor$)

Thus the running time is $O(n \lg n)$

Is this a correct asymptotic upper bound?

Is this an asymptotically tight bound?

A tighter bound is $O(n)$

How can this be? Is there a flaw in the above reasoning?



Analyzing BuildHeap(): Tight

To **Heapify** () a subtree takes $O(h)$ time where h is the height of the subtree

$h = O(\lg m)$, $m = \#$ nodes in subtree

The height of most subtrees is small

Fact: an n -element heap has at most $\lceil n/2^{h+1} \rceil$ nodes of height h



Heapsort

Given **BuildHeap()**, an in-place sorting algorithm is easily constructed:

- Maximum element is at $A[1]$

- Discard by swapping with element at $A[n]$

 - Decrement $\text{heap_size}[A]$

 - $A[n]$ now contains correct value

- Restore heap property at $A[1]$ by calling **Heapify()**

- Repeat, always swapping $A[1]$ for $A[\text{heap_size}(A)]$



Heapsort

Heapsort (A)

```
{  
    BuildHeap(A) ;  
    for (i = length(A) downto 2)  
    {  
        Swap(A[1], A[i]) ;  
        heap_size(A) -= 1 ;  
        Heapify(A, 1) ;  
    }  
}
```



Analyzing Heapsort

The call to **BuildHeap** () takes $O(n)$ time

Each of the $n - 1$ calls to **Heapify** () takes $O(\lg n)$ time

Thus the total time taken by **HeapSort** ()

$$= O(n) + (n - 1) O(\lg n)$$

$$= O(n) + O(n \lg n)$$

$$= O(n \lg n)$$



Priority Queues

Heapsort is a nice algorithm, but in practice Quicksort (coming up) usually wins

But the heap data structure is incredibly useful for implementing *priority queues*

A data structure for maintaining a set S of elements, each with an associated value or *key*

Supports the operations **Insert()**, **Maximum()**, and **ExtractMax()**

What might a priority queue be useful for?



Implementing Priority Queues with Heaps

The heap data structure with the Heapify-down and Heapify-up operations can efficiently implement a priority queue that is constrained to hold at most N elements at any point in time. Here we summarize the operations we will use.

- **StartHeap(N)** returns an empty heap H that is set up to store at most N elements. This operation takes $O(N)$ time, as it involves initializing the array that will hold the heap.
- **Insert(H, v)** inserts the item v into heap H . If the heap currently has n elements, this takes $O(\log n)$ time.
- **FindMin(H)** identifies the minimum element in the heap H but does not remove it. This takes $O(1)$ time.
- **Delete(H, i)** deletes the element in heap position i . This is implemented in $O(\log n)$ time for heaps that have n elements.
- **ExtractMin(H)** identifies and deletes an element with minimum key value from a heap. This is a combination of the preceding two operations, and so it takes $O(\log n)$ time.

